What is a windowed dataset?

- ○ The time series aligned to a fixed shape
- ○ A consistent set of subsets of a time series
- ◉ A fixed-size subset of a time series

Correct

- ○ There's no such thing

✓ 1.00/1.00 points

What does ‘drop_remainder=true’ do?

- It ensures that all data is used
- It ensures that all rows in the data window are the same length by cropping data

Correct

- It ensures that all rows in the data window are the same length by adding data
- It ensures that the data is all the same shape

✓ 1.00/1.00 points

What's the correct line of code to split an n column window into n-1 columns for features and 1 column for a label

- dataset = dataset.map(lambda window: (window[n-1], window[1]))
- dataset = dataset.map(lambda window: (window[:-1], window[-1:]))

Correct

- dataset = dataset.map(lambda window: (window[-1:], window[:-1]))
- dataset = dataset.map(lambda window: (window[n], window[1]))

✓ 1.00/1.00 points

What does MSE stand for?

- (•) Mean Squared error

Correct

- ( ) Mean Slight error
- ( ) Mean Second error
- ( ) Mean Series error

✓ 1.00/1.00 points

What does MAE stand for?

- ○ Mean Average Error
- ○ Mean Advanced Error
- ◉ Mean Absolute Error

Correct

- ○ Mean Active Error

✓ 1.00/1.00 points

If time values are in time[], series values are in series[] and we want to split the series into training and validation at time 1000, what is the correct code?

- time_train = time[:split_time]
  x_train = series[:split_time]
  time_valid = time[split_time]
  x_valid = series[split_time]

- time_train = time[split_time]
  x_train = series[split_time]
  time_valid = time[split_time]
  x_valid = series[split_time]

- time_train = time[split_time]
  x_train = series[split_time]
  time_valid = time[split_time:]
  x_valid = series[split_time:]

- time_train = time[:split_time]
  x_train = series[:split_time]
  time_valid = time[split_time:]
  x_valid = series[split_time:]

Correct

✓ 1.00/1.00 points

If you want to inspect the learned parameters in a layer after training, what's a good technique to use?

- (●) Assign a variable to the layer and add it to the model using that variable. Inspect its properties after training

Correct

- ( ) Run the model with unit data and inspect the output for that layer
- ( ) Decompile the model and inspect the parameter set for that layer
- ( ) Iterate through the layers dataset of the model to find the layer you want

✓ 1.00/1.00 points

How do you set the learning rate of the SGD optimizer?

- ○ You can’t set it
- ◉ Use the lr property

Correct

- ○ Use the RateOfLearning property
- ○ Use the Rate property

✓ 1.00/1.00 points

If you want to amend the learning rate of the optimizer on the fly, after each epoch, what do you do?

- ○ Use a LearningRateScheduler and pass it as a parameter to a callback
- ○ Callback to a custom function and change the SGD property
- ◉ Use a LearningRateScheduler object in the callbacks namespace and assign that to the callback

Correct

- ○ You can't set it